

Installing Vault



Vault is platform agnostic....meaning it can be run on many different underlying platforms, such as:



Kubernetes



Cloud-based Machines (AWS Instances, Azure Virtual Machines)



VMware Virtual Machines



Physical Servers



A Laptop

Installing Vault



Vault is also available for many operating systems...

-  macOS
-  Windows
-  Linux
-  FreeBSD
-  NetBSD
-  OpenBSD
-  Solaris

Installing Vault

Order of Operations



1 Install Vault

2 Create Configuration File

3 Initialize Vault

4 Unseal Vault

Installing Vault



▼ So where do I download Vault?

- <https://developer.hashicorp.com/vault>
- releases.hashicorp.com/vault

▼ Ideally, you should install Vault using your preferred package manager (apt, yum, homebrew, chocolately (community supported))

Terminal

```
$ sudo yum install -y yum-utils shadow-utils
$ sudo yum-config-manager --add-repo https://rpm.releases.hashicorp.com/AmazonLinux/hashicorp.repo
$ sudo yum -y install vault
```

▼ Use the Vault Helm Chart to install/configure Vault on Kubernetes

Terminal


```
$ helm install vault hashicorp/vault
```

Installing Vault


Manual Installation



Download Vault from HashiCorp



Unpackage Vault to a Directory




Set Path to Executable

Installing Vault

Package Manager



Preferred Method



Use your preferred
package manager

Running Vault Dev Server



Quickly run Vault without configuration

Automatically initialized and unsealed

Enables the UI – available at localhost

Provides an Unseal Key

Automatically logs in as root

Non-Persistent – Runs in memory

Insecure – doesn't use TLS

Sets the listener to 127.0.0.1:8200

Mounts a K/V v2 Secret Engine

Provides a root token

NEVER USE DEV SERVER MODE IN PRODUCTION!

Where Would I Use Dev Server?



Dev Server Mode



Proof of Concepts

New Development Integrations

Testing New Features of Vault

Experimenting with Features

Running Vault Dev Server



```

$ vault server -dev

==> Vault server configuration:

Administrative Namespace:
  Api Address: http://127.0.0.1:8200
  Cgo: disabled
  Cluster Address: https://127.0.0.1:8201
  Go Version: go1.23.8
  Listener 1: tcp (addr: "127.0.0.1:8200", cluster address: "127.0.0.1:8201",
disable_request_limiter: "false", max_request_duration: "1m30s", max_request_size:
"33554432", tls: "disabled")
  Log Level:
  Mlock: supported: false, enabled: false
  Recovery Mode: false
  Storage: inmem
  Version: Vault v1.19.3, built 2025-04-29T10:34:52Z
  Version Sha: a2de3bb7bcf4a073cbb8724863a5a88d3c2f83da

==> Vault server started! Log data will stream:
```



Running Vault Server in Production

- ▼ Deploy one or more persistent nodes via **configuration file**
- ▼ Use a **storage backend** that meets the requirements
- ▼ Multiple Vault nodes will be configured as a **cluster**
- ▼ Deploy close to your applications
- ▼ Most likely, you'll **automate** the provisioning of Vault



Running Vault Server in Production

- ▼ To start Vault, run the `vault server -config=<file>` command
- ▼ In a production environment, you'll have a service manager executing and managing the Vault service (systemctl, Windows Service Manager, etc)
- ▼ For Linux, you also need a `systemd` file to manage the service for Vault

Running Vault Server in Production



▼ Systemd for a Vault service:

- https://github.com/btkrausen/hashicorp/blob/master/vault/config_files/vault.service

▼ Systemd file for a Consul Server:

- <https://github.com/btkrausen/hashicorp/blob/master/consul/consul.service>


▼ Systemd for a Consul client (that would run on the Vault node):

- https://github.com/btkrausen/hashicorp/blob/master/vault/config_files/consul-client.json



Running Vault Server in Production

Single Node




Not a Recommended Architecture

- ▼ No Redundancy
- ▼ No Scalability

Running Vault Server in Production




Multi-Node Vault Cluster (with Integrated Storage)





Running Vault Server in Production

Multi-Node Vault Cluster (with External Storage)



Deploying the Integrated Storage Backend



Vault Internal Storage Option

Leverages Raft Consensus Protocol

All Vault nodes have a copy of the data

Eliminates Network Hop to Consul

Supports High Availability

Only need to troubleshoot Vault

Built-in Snapshots For Data Retention

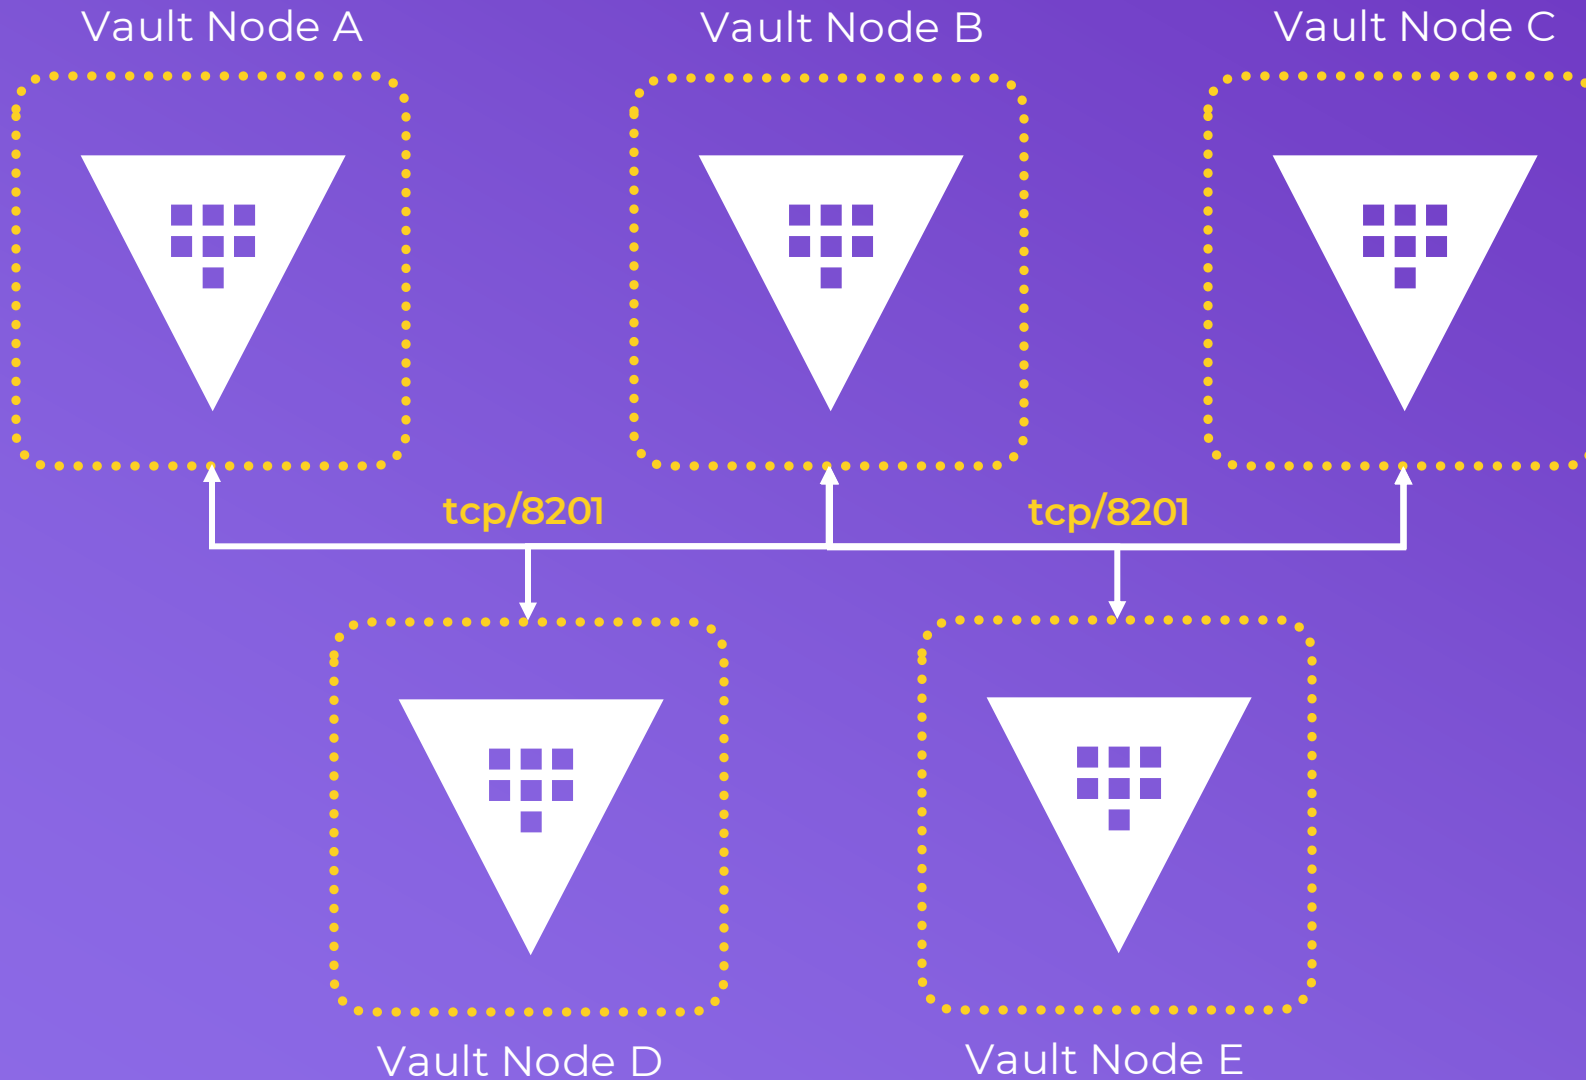
HashiCorp Supported

Deploying the Integrated Storage Backend



- ▼ Integrated Storage (aka **Raft**) allows Vault nodes to provide its own replicated storage across the Vault nodes within a cluster
- ▼ Define a **local path** to store replicated data
- ▼ All data is replicated among **all nodes** in the cluster
- ▼ Eliminates the need to run a Consul cluster and manage it

Deploying the Integrated Storage Backend



Deploying the Integrated Storage Backend



Example Vault Server Configuration File

```
storage "raft" {
  path      = "/opt/vault/data"
  node_id   = "node-a-us-east-1.example.com"
  retry_join {
    auto_join = "provider=aws region=us-east-1 tag_key=vault tag_value=us-east-1"
  }
}


listener "tcp" {
  address       = "0.0.0.0:8200"
  cluster_address = "0.0.0.0:8201"
  tls_disable   = 0
  tls_cert_file = "/etc/vault.d/client.pem"
  ...
}
```

Deploying the Integrated Storage Backend



Manually join standby nodes to the cluster using the CLI:

```
TERMINAL
$ vault operator raft join https://active_node.example.com:8200
```



Deploying the Integrated Storage Backend



List the cluster members

```

$ vault operator raft list-peers

Node      Address          State    Voter
----      -
vault_1   10.0.101.22:8201 leader    true
vault_2   10.0.101.23:8201 follower  true
vault_3   10.0.101.24:8201 follower  true
vault_4   10.0.101.25:8201 follower  true
vault_5   10.0.101.26:8201 follower  true

```